

Programming elements:

Classification

Assignment:

1. (Titanic Dataset)

1. Find the correlation between 'survived' (target column) and 'sex' column for the Titanic use case in class.
 - a. Do you think we should keep this feature?
2. Do at least two visualizations to describe or show correlations.
3. Implement Naïve Bayes method using scikit-learn library and report the accuracy.

Question 1

```
In [2]: import pandas as pd
import numpy as np
import random as rnd

import warnings # current version generates a bunch of warnings that we'll ignore
warnings.filterwarnings("ignore")

# visualization
import seaborn as sns
import matplotlib.pyplot as plt

# machine Learning
from sklearn.linear_model import LogisticRegression, RidgeClassifierCV
from sklearn.svm import SVC, LinearSVC
from sklearn.ensemble import (RandomForestClassifier, GradientBoostingClassifier)
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.metrics import accuracy_score
from sklearn import preprocessing
```

First of all, imported pandas as pd ,numpy as np,random as rnd, warnings, seaborn as sns, matplotlib.pyplot as plt and import seaborn as sns

from sklearn.linear_model import LogisticRegression, RidgeClassifierCV

from sklearn.svm import SVC, LinearSVC

from sklearn.ensemble import (RandomForestClassifier, GradientBoostingClassifier)

from sklearn.neighbors import KNeighborsClassifier

from sklearn.naive_bayes import GaussianNB

from sklearn.linear_model import Perceptron

from sklearn.linear_model import SGDClassifier

from sklearn.tree import DecisionTreeClassifier

from sklearn.model_selection import cross_val_score, GridSearchCV

```
from sklearn.metrics import accuracy_score
```

```
from sklearn import preprocessing
```

#Question-1 (TITANIC DATASET)

```
train_df = pd.read_csv('train.csv')
```

```
train_df.head()
```

The datasets are read using read_csv() and Sex and Embark are replaced with numerical values for finding the correlation heatmap shows the correlation and Sex has the highest value of the correlation.

```
In [5]: train_df = pd.read_csv('train.csv')
train_df.head()
```

Out[5]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cummings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
j]: #1.. correlation between 'survived' (target column) and 'sex' column for the Titanic use case in class.
# a. Why should we keep this feature?

#print(pd.pivot_table(train_df, index = 'Survived', columns = 'Sex',
# values = 'Ticket',aggfunc = 'count'))

train_df[["Sex", "Survived"]].groupby(['Sex'], as_index=False).mean().sort_values(by='Survived', ascending=False)
#a. Survived: Most of the people died, only around 300 people survived.
# Sex: There were more males than females aboard the ship, roughly double the amount.
#Most of the women survived, and the majority of the male died .

...+rct.
```

#1.. correlation between 'survived' (target column) and 'sex' column for the Titanic use case in class.

a. Why should we keep this feature?

```
#print(pd.pivot_table(train_df, index = 'Survived', columns = 'Sex',
# values = 'Ticket',aggfunc = 'count'))
```

```
train_df[["Sex", "Survived"]].groupby(['Sex'], as_index=False).mean().sort_values(by='Survived',
ascending=False)
```

#a. Survived: Most of the people died, only around 300 people survived.

Sex: There were more males than females aboard the ship, roughly double the amount.

#Most of the women survived, and the majority of the male died .

Output will be as

Out[6]:

	Sex	Survived
0	female	0.742038
1	male	0.188908

```
In [7]: a = preprocessing.LabelEncoder()
train_df['Sex'] = a.fit_transform(train_df.Sex.values)
train_df['Survived'].corr(train_df['Sex'])
```

Output is -0.5433513806577552

```
In [8]: mat = train_df.corr()
print(mat)
```

mat = train_df.corr()

print(mat)

Output for the above code is

	PassengerId	Survived	Pclass	Sex	Age	SibSp	\
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.036847	-0.057527	
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	
Age	0.036847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	
Parch	-0.001652	0.081629	0.018443	-0.245489	-0.189119	0.414838	
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159651	

	Parch	Fare
PassengerId	-0.001652	0.012658
Survived	0.081629	0.257307
Pclass	0.018443	-0.549500
Sex	-0.245489	-0.182333
Age	-0.189119	0.096067
SibSp	0.414838	0.159651
Parch	1.000000	0.216225
Fare	0.216225	1.000000

#2.. two visualizations to describe or show correlations

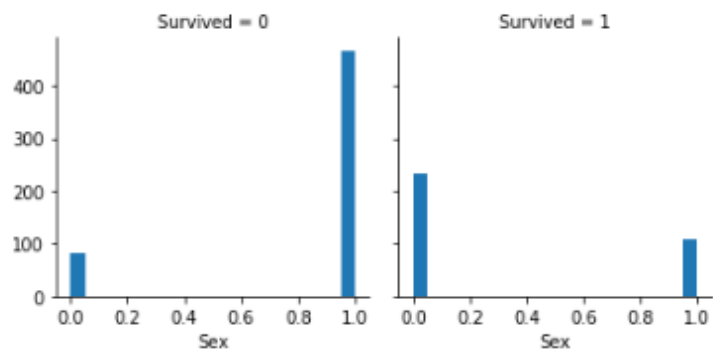
```
g = sns.FacetGrid(train_df, col='Survived')
```

```
g.map(plt.hist, 'Sex', bins=20)
```

```
In [9]: #2.. two visualizations to describe or show correlations
g = sns.FacetGrid(train_df, col='Survived')
g.map(plt.hist, 'Sex', bins=20)
```

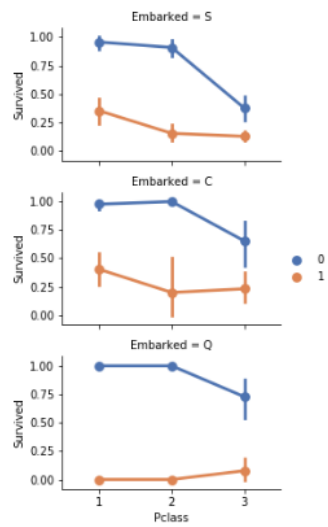
```
Out[9]: <seaborn.axisgrid.FacetGrid at 0x2147b49cc48>
```

```
Out[9]: <seaborn.axisgrid.FacetGrid at 0x2147b49cc48>
```



```
In [10]: grid = sns.FacetGrid(train_df, row='Embarked', size=2.2, aspect=1.6)
grid.map(sns.pointplot, 'Pclass', 'Survived', 'Sex', palette='deep')
grid.add_legend()
```

```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x2147c170c48>
```




```
In [11]: train_df.corr().style.background_gradient(cmap="Greens")
```

Out[11]:

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
PassengerId	1.000000	-0.005007	-0.035144	0.042939	0.038847	-0.057527	-0.001852	0.012658
Survived	-0.005007	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307
Pclass	-0.035144	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500
Sex	0.042939	-0.543351	0.131900	1.000000	0.093254	-0.114631	-0.245489	-0.182333
Age	0.038847	-0.077221	-0.369226	0.093254	1.000000	-0.308247	-0.189119	0.096067
SibSp	-0.057527	-0.035322	0.083081	-0.114631	-0.308247	1.000000	0.414838	0.159051
Parch	-0.001852	0.081629	0.018443	-0.245489	-0.189119	0.414838	1.000000	0.216225
Fare	0.012658	0.257307	-0.549500	-0.182333	0.096067	0.159051	0.216225	1.000000

```
In [12]: sns.heatmap(mat, annot=True, vmax=1, vmin=-1, center=0, cmap='vlag')
plt.show()
```



```
In [13]: #3.. Implementing Naïve Bayes method using scikit-Learn library and report the accuracy
```

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
```

```
In [15]: train_raw = pd.read_csv('train.csv')
test_raw = pd.read_csv('test.csv')

# Join data to analyse and process the set as one.
train_raw['train'] = 1
test_raw['train'] = 0
df = train_raw.append(test_raw, sort=False)

features = ['Age', 'Embarked', 'Fare', 'Parch', 'Pclass', 'Sex', 'SibSp']
target = 'Survived'

df = df[features + [target] + ['train']]
# Categorical values need to be transformed into numeric.
df['Sex'] = df['Sex'].replace(["female", "male"], [0, 1])
df['Embarked'] = df['Embarked'].replace(['S', 'C', 'Q'], [1, 2, 3])
train = df.query('train == 1')
test = df.query('train == 0')
```

```
In [16]: # Drop missing values from the train set.
train.dropna(axis=0, inplace=True)
labels = train[target].values
train.drop(['train', target, 'Pclass'], axis=1, inplace=True)
test.drop(['train', target, 'Pclass'], axis=1, inplace=True)
from sklearn.model_selection import train_test_split, cross_validate

X_train, X_test, y_train, y_test = train_test_split(train, labels, test_size=0.2, random_state=0)
```

```
In [17]: # Gaussian Naive Bayes
from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X_train, y_train)

y_pred = classifier.predict(X_test)

# Summary of the predictions made by the classifier
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(y_pred, y_test))
```

	precision	recall	f1-score	support
0.0	0.77	0.84	0.80	82
1.0	0.75	0.66	0.70	61
accuracy			0.76	143
macro avg	0.76	0.75	0.75	143
weighted avg	0.76	0.76	0.76	143

```
[[69 13]
 [21 40]]
accuracy is 0.7622377622377622
```

#Question-2 (GLASS DATASET)

#1.. Implementing Naïve Bayes method using scikit-learn library.

a. Use the glass dataset available

Train_test_split() is used to split the training and testing data and random state=0 gives the same testing and training split every time.

```
In [19]: #1.. Implementing Naïve Bayes method using scikit-learn library.
# a. Use the glass dataset available

glass = pd.read_csv('glass.csv')
glass.head()
```

```
Out[19]:
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
1	1.51761	13.89	3.80	1.36	72.73	0.48	7.83	0.0	0.0	1
2	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
3	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1
4	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.0	0.0	1

```
In [20]: glass.corr().style.background_gradient(cmap="Greens")
```

```
Out[20]:
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
RI	1.000000	-0.191885	-0.122274	-0.407326	-0.542052	-0.289833	0.810403	-0.000386	0.143010	-0.164237
Na	-0.191885	1.000000	-0.273732	0.156794	-0.069809	-0.268087	-0.275442	0.326803	-0.241346	0.502898
Mg	-0.122274	-0.273732	1.000000	-0.481799	-0.165927	0.005396	-0.443750	-0.492262	0.083060	-0.744993
Al	-0.407326	0.156794	-0.481799	1.000000	-0.005524	0.325958	-0.259592	0.479404	-0.074402	0.598829
Si	-0.542052	-0.069809	-0.165927	-0.005524	1.000000	-0.193331	-0.208732	-0.102151	-0.094201	0.151565
K	-0.289833	-0.268087	0.005396	0.325958	-0.193331	1.000000	-0.317836	-0.042618	-0.007719	-0.010054
Ca	0.810403	-0.275442	-0.443750	-0.259592	-0.208732	-0.317836	1.000000	-0.112841	0.124968	0.000952
Ba	-0.000386	0.326803	-0.492262	0.479404	-0.102151	-0.042618	-0.112841	1.000000	-0.058692	0.575161
Fe	0.143010	-0.241346	0.083060	-0.074402	-0.094201	-0.007719	0.124968	-0.058692	1.000000	-0.188278
Type	-0.164237	0.502898	-0.744993	0.598829	0.151565	-0.010054	0.000952	0.575161	-0.188278	1.000000

```
In [21]: y=glass['Type'].values
```

```
In [22]: #1b. Use train_test_split to create training and testing part.
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size = 0.30, random_state = 0)
```

#1b. Use train_test_split to create training and testing part.

```
In [23]: # Evaluating the model on testing part using score and
# 1. Gaussian Naive Bayes
from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(x_train, y_train)

y_pred = classifier.predict(x_test)

# Summary of the predictions made by the classifier
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is',accuracy_score(y_pred,y_test))
```

	precision	recall	f1-score	support
1	0.39	0.86	0.54	21
2	0.50	0.12	0.19	26
3	0.00	0.00	0.00	7
5	0.00	0.00	0.00	2
6	0.67	1.00	0.80	2
7	0.88	1.00	0.93	7
accuracy			0.46	65
macro avg	0.41	0.50	0.41	65
weighted avg	0.44	0.46	0.37	65

[18	1	0	0	1	1]
[21	3	1	1	0	0]
[7	0	0	0	0	0]
[0	2	0	0	0	0]
[0	0	0	0	2	0]
[0	0	0	0	0	7]]

accuracy is 0.46153846153846156

Multinomial Naive Bayes and Gaussian Naive Bayes gives accurate values in Naive Bayes

#Question-3

```
In [24]: #1. Implement Linear SVM method using scikit library
# a. Use the glass dataset available
# Support Vector Machine's
from sklearn.svm import SVC, LinearSVC

classifier = LinearSVC()
classifier.fit(x_train, y_train)

y_pred = classifier.predict(x_test)

# Summary of the predictions made by the classifier
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
# Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(y_pred, y_test))
```

	precision	recall	f1-score	support
1	0.43	1.00	0.60	21
2	0.50	0.12	0.19	26
3	0.00	0.00	0.00	7
5	0.00	0.00	0.00	2
6	0.00	0.00	0.00	2
7	0.78	1.00	0.88	7
accuracy			0.48	65
macro avg	0.28	0.35	0.28	65
weighted avg	0.42	0.48	0.36	65

```
[[21 0 0 0 0 0]
 [21 3 0 0 1 1]
 [ 7 0 0 0 0 0]
 [ 0 2 0 0 0 0]
 [ 0 1 0 0 0 1]
 [ 0 0 0 0 0 7]]
accuracy is 0.47692307692307695
```

#1. Implement linear SVM method using scikit library

a. Use the glass dataset available

Support Vector Machine's

from sklearn.svm import SVC, LinearSVC

classifier = LinearSVC()

classifier.fit(x_train, y_train)

y_pred = classifier.predict(x_test)

Summary of the predictions made by the classifier

print(classification_report(y_test, y_pred))

print(confusion_matrix(y_test, y_pred))

Accuracy score



```
from sklearn.metrics import accuracy_score

print('accuracy is',accuracy_score(y_pred,y_test))
```

#Do at least two visualizations to describe or show correlations in the Glass Dataset

```
In [25]: #Do at least two visualizations to describe or show correlations in the Glass Dataset
g = sns.FacetGrid(glass, col='Type')
g.map(plt.hist, 'RI', bins=20)
```

Out[25]: <seaborn.axisgrid.FacetGrid at 0x2147c6b92c8>



```
In [26]: grid = sns.FacetGrid(glass, row='Type', col='Ba', height=2.2, aspect=1.6)
grid.map(sns.barplot, 'Al', 'Ca', alpha=.5, ci=None)
grid.add_legend()
```

Out[26]: <seaborn.axisgrid.FacetGrid at 0x2147caff288>



```
In [ ]: #Which algorithm you got better accuracy? Can you justify why?
#Gaussian Naive Bayes algorithm gives better accuracy than other algorithms. This is used when features are not discreet.
```

GitHUB Link: https://github.com/AishaFar/ML_Assignment3_Farhana_700735341/

Video Link:

https://youtu.be/mY6EB_6wdt8